

Chapter 3

The Relational Model

1. The first database systems were based on the **network** and **hierarchical** models. These are covered briefly in appendices in the text. The relational model was first proposed by E.F. Codd in 1970 and the first such systems (notably INGRES and System/R) was developed in 1970s. The relational model is now the dominant model for commercial data processing applications.
2. Note: **Attribute Name Abbreviations**

The text uses fairly long attribute names which are abbreviated in the notes as follows.

- o *customer-name* becomes *cname*
 - o *customer-city* becomes *ccity*
 - o *branch-city* becomes *bcity*
 - o *branch-name* becomes *bname*
 - o *account-number* becomes *account#*
 - o *loan-number* becomes *loan#*
 - o *banker-name* becomes *banker*
- [Structure of Relational Database](#)
 - o [Basic Structure](#)
 - o [Database Scheme](#)
 - o [Keys](#)
 - o [Query Languages](#)
- [The Relational Algebra](#)
 - o [Fundamental Operations](#)
 - o [Formal Definition of Relational Algebra](#)
 - o [Additional Operations](#)
- [The Tuple Relational Calculus](#)
 - o [Example Queries](#)
 - o [Formal Definitions](#)
 - o [Safety of Expressions](#)
 - o [Expressive Power of Languages](#)
- [The Domain Relational Calculus](#)
 - o [Formal Definitions](#)
 - o [Example Queries](#)
 - o [Safety of Expressions](#)
 - o [Expressive Power of Languages](#)
- [Modifying the Database](#)

- o [Deletion](#)
- o [Insertions](#)
- o [Updating](#)
- [Views](#)
 - o [View Definition](#)
 - o [Updates Through Views and Null Values](#)

Structure of Relational Database

1. A relational database consists of a collection of **tables**, each having a unique **name**.

A **row** in a table represents a **relationship** among a set of values.

Thus a table represents a **collection of relationships**.

2. There is a direct correspondence between the concept of a table and the mathematical concept of a relation. A substantial theory has been developed for relational databases.

Basic Structure

1. Figure 3.1 shows the *deposit* and *customer* tables for our banking example.

bname	account#	cname	balance
Downtown	101	Johnson	500
Longhead Mall	215	Smith	700
SFU	102	Hayes	400
SFU	301	Adams	1300

cname	street	city
Johnson	Pender	Vancouver
Smith	North	Burnaby
Hayes	Curtis	Burnaby
Adams	No 3 Road	Richmond
Jones	Oak	Vancouver

Figure 3.1: The *deposit* and *customer* relations.

- o It has four attributes.
- o For each attribute there is a permitted set of values, called the **domain** of that attribute.
- o E.g. the domain of *bname* is the set of all branch names.

Let D_1 denote the domain of *bname*, and D_2 , D_3 and D_4 the remaining attributes' domains respectively.

Then, any row of *deposit* consists of a four-tuple (v_1, v_2, v_3, v_4) where

$$v_1 \in D_1, v_2 \in D_2, v_3 \in D_3, v_4 \in D_4$$

In general, *deposit* contains a **subset** of the set of all possible rows.

That is, *deposit* is a subset of

$$D_1 \times D_2 \times D_3 \times D_4, \quad \text{or, abbreviated to,} \quad \times_{i=1}^4 D_i$$

In general, a table of n columns must be a subset of

$$\times_{i=1}^n D_i \quad (\text{all possible rows})$$

2. Mathematicians define a relation to be a subset of a Cartesian product of a list of domains. You can see the correspondence with our tables.

We will use the terms **relation** and **tuple** in place of **table** and **row** from now on.

3. Some more formalities:
 - o let the tuple variable t refer to a tuple of the relation r .
 - o We say $t \in r$ to denote that the tuple t is in relation r .
 - o Then $t[bname] = t[1] =$ the value of t on the *bname* attribute.
 - o So $t[bname] = t[1] = \text{"Downtown"}$,
 - o and $t[cname] = t[3] = \text{"Johnson"}$.
4. We'll also require that the domains of all attributes be **indivisible units**.
 - o A domain is **atomic** if its elements are indivisible units.
 - o For example, the set of integers is an atomic domain.
 - o The set of all sets of integers is not.
 - o Why? Integers do not have subparts, but sets do - the integers comprising them.
 - o We could consider integers non-atomic if we thought of them as ordered lists of digits.

Database Scheme

1. We distinguish between a **database scheme** (logical design) and a **database instance** (data in the database at a point in time).
2. A **relation scheme** is a list of attributes and their corresponding domains.
3. The text uses the following conventions:
 - o italics for all names
 - o lowercase names for relations and attributes
 - o names beginning with an uppercase for relation schemes

These notes will do the same.

For example, the relation scheme for the *deposit* relation:

- o $Deposit\text{-}scheme = (bname, account\#, cname, balance)$

We may state that *deposit* is a relation on scheme *Deposit-scheme* by writing *deposit(Deposit-scheme)*.

If we wish to specify domains, we can write:

- o $(bname: \text{string}, account\#: \text{integer}, cname: \text{string}, balance: \text{integer})$.

Note that customers are identified by name. In the real world, this would not be allowed, as two or more customers might share the same name.

Figure 3.2 shows the E-R diagram for a banking enterprise.

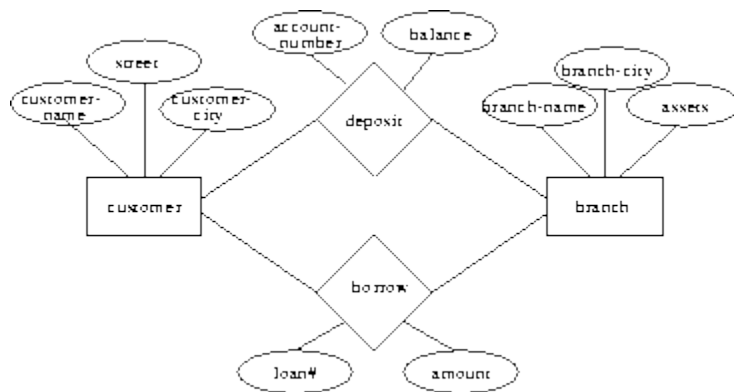


Figure 3.2: E-R diagram for the banking enterprise

- The relation schemes for the banking example used throughout the text are:
 - o $Branch\text{-}scheme = (bname, assets, bcity)$
 - o $Customer\text{-}scheme = (cname, street, ccity)$
 - o $Deposit\text{-}scheme = (bname, account\#, cname, balance)$
 - o $Borrow\text{-}scheme = (bname, loan\#, cname, amount)$

Note: some attributes appear in several relation schemes (e.g. *bname*, *cname*). This is legal, and provides a way of **relating** tuples of distinct relations.

5. Why not put all attributes in one relation?

Suppose we use one large relation instead of *customer* and *deposit*:

- o $Account\text{-}scheme = (bname, account\#, cname, balance, street, ccity)$
- o If a customer has several accounts, we must duplicate her or his address for each account.
- o If a customer has an account but no current address, we cannot build a tuple, as we have no values for the address.
- o We would have to use **null values** for these fields.

- o Null values cause difficulties in the database.
- o By using two separate relations, we can do this without using null values

Keys

1. The notions of **superkey**, **candidate key** and **primary key** all apply to the relational model.
2. For example, in *Branch-scheme*,
 - o $\{bname\}$ is a superkey.
 - o $\{bname, bcity\}$ is a superkey.
 - o $\{bname, bcity\}$ is not a candidate key, as the superkey $\{bname\}$ is contained in it.
 - o $\{bname\}$ is a candidate key.
 - o $\{bcity\}$ is not a superkey, as branches may be in the same city.
 - o We will use $\{bname\}$ as our primary key.
3. The primary key for *Customer-scheme* is $\{cname\}$.
4. More formally, if we say that a subset K of R is a *superkey* for R , we are restricting consideration to relations $r(R)$ in which no two distinct tuples have the same values on all attributes in K . In other words,
 - o If t_1 and t_2 are in r , and
 - o $t_1 \neq t_2$,
 - o then $t_1[K] \neq t_2[K]$.

Query Languages

1. A **query language** is a language in which a user requests information from a database. These are typically higher-level than programming languages.

They may be one of:

- o **Procedural**, where the user instructs the system to perform a sequence of operations on the database. This will compute the desired information.
 - o **Nonprocedural**, where the user specifies the information desired without giving a procedure for obtaining the information.
2. A complete query language also contains facilities to insert and delete tuples as well as to modify parts of existing tuples.

The Relational Algebra

1. The **relational algebra** is a procedural query language.

- o Six fundamental operations:
 - select (unary)
 - project (unary)
 - rename (unary)
 - cartesian product (binary)
 - union (binary)
 - set-difference (binary)
- o Several other operations, defined in terms of the fundamental operations:
 - set-intersection
 - natural join
 - division
 - assignment
- o Operations produce a new relation as a re

Fundamental Operations

1. The Select Operation

Select selects tuples that satisfy a given predicate. Select is denoted by a lowercase Greek sigma (σ), with the predicate appearing as a subscript. The argument relation is given in parentheses following the σ .

For example, to select tuples (rows) of the *borrow* relation where the branch is "SFU", we would write

$$\sigma_{branch = "SFU"}(borrow)$$

Let Figure 3.3 be the *borrow* and *branch* relations in the banking example.

bname	loan#	cname	amount
Downtown	17	Jones	1000
Longheed.Mall	23	Smith	2000
SFU	15	Hayes	1500

bname	assets	city
Downtown	9,000,000	Vancouver
Longheed.Mall	21,000,000	Burnaby
SFU	17,000,000	Burnaby

Figure 3.3: The *borrow* and *branch* relations.

The new relation created as the result of this operation consists of one tuple: (*SFU*, 15, *Hayes*, 1500).

We allow comparisons using =, \neq , <, \leq , > and \geq in the selection predicate.

We also allow the logical connectives \vee (or) and \wedge (and). For example:

$$\sigma_{branch = "Downtown" \wedge amount > 1200}(borrow)$$

cname	banker
Hayes	Jones
Johnson	Johnson

Figure 3.4: The *client* relation.

Suppose there is one more relation, *client*, shown in Figure 3.4, with the scheme

$$Client_scheme = (cname, banker)$$

we might write

$$\sigma_{cname=banker}(client)$$

to find clients who have the same name as their banker.

2. The Project Operation

Project copies its argument relation for the specified attributes only. Since a relation is a **set**, duplicate rows are eliminated.

Projection is denoted by the Greek capital letter pi (Π). The attributes to be copied appear as subscripts.

For example, to obtain a relation showing customers and branches, but ignoring amount and loan#, we write

$$\Pi_{branch, cname}(borrow)$$

We can perform these operations on the relations resulting from other operations.

To get the names of customers having the same name as their bankers,

$$\Pi_{cname}(\sigma_{cname=banker}(client))$$

Think of **select** as taking rows of a relation, and **project** as taking columns of a relation.

3. The Cartesian Product Operation

The **cartesian product** of two relations is denoted by a cross (\times), written

$$r_1 \times r_2 \text{ for relations } r_1 \text{ and } r_2$$

The result of $r_1 \times r_2$ is a new relation with a tuple for each possible **pairing** of tuples from r_1 and r_2 .

In order to avoid ambiguity, the attribute names have attached to them the name of the relation from which they came. If no ambiguity will result, we drop the relation name.

The result $client \times customer$ is a very large relation. If r_1 has n_1 tuples, and r_2 has n_2 tuples, then $r = r_1 \times r_2$ will have $n_1 n_2$ tuples.

The resulting scheme is the concatenation of the schemes of r_1 and r_2 , with relation names added as mentioned.

To find the clients of banker Johnson and the city in which they live, we need information in both *client* and *customer* relations. We can get this by writing

$$\sigma_{\text{banker} = \text{"Johnson"}}(client \times customer)$$

However, the *customer.cname* column contains customers of bankers other than Johnson. (Why?)

We want rows where *client.cname* = *customer.cname*. So we can write

$$\sigma_{\text{client.cname} = \text{customer.cname}}(\sigma_{\text{banker} = \text{"Johnson"}}(client \times customer))$$

to get just these tuples.

Finally, to get just the customer's name and city, we need a projection:

$$\Pi_{\text{client.cname, city}}(\sigma_{\text{client.cname} = \text{customer.cname}}(\sigma_{\text{banker} = \text{"Johnson"}}(client \times customer)))$$

4. The Rename Operation

The **rename** operation solves the problems that occurs with naming when performing the cartesian product of a relation with itself.

Suppose we want to find the names of all the customers who live on the same street and in the same city as Smith.

We can get the street and city of Smith by writing

$$\Pi_{\text{street, city}}(\sigma_{\text{cname} = \text{"Smith"}}(customer))$$

To find other customers with the same information, we need to reference the *customer* relation again:

$$\sigma_P(customer \times (\Pi_{\text{street, city}}(\sigma_{\text{cname} = \text{"Smith"}}(customer))))$$

where P is a selection predicate requiring *street* and *city* values to be equal.

Problem: how do we distinguish between the two street values appearing in the Cartesian product, as both come from a *customer* relation?

Solution: use the rename operator, denoted by the Greek letter rho (ρ).

We write

$$\rho_x(r)$$

to get the relation r under the name of x .

If we use this to rename one of the two **customer** relations we are using, the ambiguities will disappear.

$$\Pi_{customer.cname}(\sigma_{cust2.street=customer.street \wedge cust2.city=customer.city} (customer \times (\Pi_{street,city}(\sigma_{cname="Smith"}(\rho_{cust2}(customer))))))$$

5. The Union Operation

The **union** operation is denoted \cup as in set theory. It returns the union (set union) of two compatible relations.

For a union operation $r \cup s$ to be legal, we require that

- o r and s must have the same number of attributes.
- o The domains of the corresponding attributes must be the same.

To find all customers of the SFU branch, we must find everyone who has a loan or an account or both at the branch.

We need both *borrow* and *deposit* relations for this:

$$\Pi_{cname}(\sigma_{branch="SFU"}(borrow)) \cup \Pi_{cname}(\sigma_{branch="SFU"}(deposit))$$

As in all set operations, duplicates are eliminated, giving the relation of Figure 3.5(a).

(a)	<table><tr><th>cname</th></tr><tr><td>Hayes</td></tr><tr><td>Adams</td></tr></table>	cname	Hayes	Adams	(b)	<table><tr><th>cname</th></tr><tr><td>Adams</td></tr></table>	cname	Adams
cname								
Hayes								
Adams								
cname								
Adams								

Figure 3.5: The *union* and *set-difference* operations.

6. The Set Difference Operation

Set difference is denoted by the minus sign ($-$). It finds tuples that are in one relation, but not in another.

Thus $r - s$ results in a relation containing tuples that are in r but not in s .

To find customers of the SFU branch who have an account there but no loan, we write

$$\Pi_{\text{customer}}(\sigma_{\text{branch} = \text{"SFU"}}(\text{deposit})) - \Pi_{\text{customer}}(\sigma_{\text{branch} = \text{"SFU"}}(\text{borrow}))$$

The result is shown in Figure 3.5(b).

We can do more with this operation. Suppose we want to find the largest account balance in the bank.

Strategy:

- o Find a relation r containing the balances **not** the largest.
- o Compute the set difference of r and the *deposit* relation.

To find r , we write

$$\Pi_{\text{deposit.balance}}(\sigma_{\text{deposit.balance} < d.\text{balance}}(\text{deposit} \times \rho_d(\text{deposit})))$$

This resulting relation contains all balances except the largest one. (See Figure 3.6(a)).

Now we can finish our query by taking the set difference:

$$\Pi_{\text{balance}}(\text{deposit}) - \Pi_{\text{deposit.balance}}(\sigma_{\text{deposit.balance} < d.\text{balance}}(\text{deposit} \times \rho_d(\text{deposit})))$$

Figure 3.6(b) shows the result.

(a)

balance
400
500
700

(b)

balance
1300

Figure 3.6: Find the largest account balance in the bank.

Formal Definition of Relational Algebra

1. A basic expression consists of either
 - o A relation in the database.
 - o A constant relation.
2. General expressions are formed out of smaller subexpressions using
 - o $\sigma_p(E_1)$ select (p a predicate)
 - o $\Pi_s(E_1)$ project (s a list of attributes)
 - o $\rho_x(E_1)$ rename (x a relation name)
 - o $E_1 \cup E_2$ union
 - o $E_1 - E_2$ set difference

- o $E_1 \times E_2$ cartesian product

Additional Operations

1. Additional operations are defined in terms of the fundamental operations. They do not add power to the algebra, but are useful to simplify common queries.
2. **The Set Intersection Operation**

Set intersection is denoted by \cap , and returns a relation that contains tuples that are in **both** of its argument relations.

It does not add any power as

$$r \cap s = r - (r - s)$$

To find all customers having both a loan and an account at the SFU branch, we write

$$\Pi_{cname}(\sigma_{branch='SFU'}(borrow)) \cap \Pi_{cname}(\sigma_{branch='SFU'}(deposit))$$

3. **The Natural Join Operation**

Often we want to simplify queries on a cartesian product.

For example, to find all customers having a loan at the bank and the cities in which they live, we need *borrow* and *customer* relations:

$$\Pi_{borrow.cname,city}(\sigma_{borrow.cname=customer.cname}(borrow \times customer))$$

Our selection predicate obtains only those tuples pertaining to only one *cname*.

This type of operation is very common, so we have the **natural join**, denoted by a \bowtie sign. Natural join combines a cartesian product and a selection into one operation. It performs a selection forcing equality on those attributes that appear in both relation schemes. Duplicates are removed as in all relation operations.

To illustrate, we can rewrite the previous query as

$$\Pi_{cname,city}(borrow \bowtie customer)$$

The resulting relation is shown in Figure [3.7](#).

cname	city
Smith	Burnaby
Hayes	Burnaby
Jones	Vancouver

Figure 3.7: Joining *borrow* and *customer* relations.

We can now make a more formal definition of natural join.

- o Consider R and S to be **sets** of attributes.
- o We denote attributes appearing in **both** relations by $R \cap S$.
- o We denote attributes in **either or both** relations by $R \cup S$.
- o Consider two relations $r(R)$ and $s(S)$.
- o The natural join of r and s , denoted by $r \bowtie s$ is a relation on scheme $R \cup S$.
- o It is a projection onto $R \cup S$ of a selection on $r \times s$ where the predicate requires $r.A = s.A$ for each attribute A in $R \cap S$.

Formally,

$$r \bowtie s = \Pi_{R \cup S} (\sigma_{r.A_1 = s.A_1 \wedge r.A_2 = s.A_2 \wedge \dots \wedge r.A_n = s.A_n} (r \times s))$$

where $R \cap S = \{A_1, A_2, \dots, A_n\}$.

To find the assets and names of all branches which have depositors living in Stamford, we need *customer*, *deposit* and *branch* relations:

$$\Pi_{branch, assets} (\sigma_{city = \text{"Stamford"}} (customer \bowtie deposit \bowtie branch))$$

Note that \bowtie is associative.

To find all customers who have both an account and a loan at the SFU branch:

$$\Pi_{names} (\sigma_{branch = \text{"SFU"}} (borrow \bowtie deposit))$$

This is equivalent to the set intersection version we wrote earlier. We see now that there can be several ways to write a query in the relational algebra.

If two relations $r(R)$ and $s(S)$ have no attributes in common, then $R \cap S = \emptyset$, and $r \bowtie s = r \times s$.

4. The Division Operation

Division, denoted \div , is suited to queries that include the phrase "for all".

Suppose we want to find all the customers who have an account at **all** branches located in Brooklyn.

Strategy: think of it as three steps.

We can obtain the names of all branches located in Brooklyn by

$$r_1 = \Pi_{names} (\sigma_{city = \text{"Brooklyn"}} (branch))$$

Figure 3.19 in the textbook shows the result.

We can also find all *cname*, *bname* pairs for which the customer has an account by

$$r_2 = \Pi_{cname, bname}(deposit)$$

Figure 3.20 in the textbook shows the result.

Now we need to find all customers who appear in r_2 with **every** branch name in r_1 .

The divide operation provides exactly those customers:

$$\Pi_{cname, bname}(deposit) \div \Pi_{bname}(\sigma_{city = 'Brooklyn'}(branch))$$

which is simply $r_2 \div r_1$.

Formally,

- o Let $r(R)$ and $s(S)$ be relations.
- o Let $S \subseteq R$.
- o The relation $r \div s$ is a relation on scheme $R - S$.
- o A tuple t is in $r \div s$ if for every tuple t_s in s there is a tuple t_r in r satisfying both of the following:

$$t_r[S] = t_s[S] \tag{3.2.1}$$

$$t_r[R - S] = t[R - S] \tag{3.2.2}$$

- o These conditions say that the $R - S$ portion of a tuple t is in $r \div s$ if and only if there are tuples with the $r - s$ portion **and** the S portion in r for **every** value of the S portion in relation S .

We will look at this explanation in class more closely.

The division operation can be defined in terms of the fundamental operations.

$$r \div s = \Pi_{R-S}(r) - \Pi_{R-S}((\Pi_{R-S}(r) \times s) - r)$$

Read the text for a more detailed explanation.

5. The Assignment Operation

Sometimes it is useful to be able to write a relational algebra expression in parts using a temporary relation variable (as we did with r_1 and r_2 in the division example).

The assignment operation, denoted \leftarrow , works like assignment in a programming language.

We could rewrite our division definition as

$$\begin{aligned}temp &\leftarrow \Pi_{R-S}(r) \\ temp &- \Pi_{R-S}((temp \times s) - r)\end{aligned}$$

No extra relation is added to the database, but the relation variable created can be used in subsequent expressions. Assignment to a permanent relation would constitute a modification to the database.

The Tuple Relational Calculus

1. The tuple relational calculus is a nonprocedural language. (The relational algebra was procedural.)

We must provide a formal description of the information desired.

2. A query in the tuple relational calculus is expressed as

$$\{t \mid P(t)\}$$

i.e. the set of tuples t for which predicate P is true.

3. We also use the notation
 - o $t[a]$ to indicate the value of tuple t on attribute a .
 - o $t \in r$ to show that tuple t is in relation r .

Example Queries

1. For example, to find the branch-name, loan number, customer name and amount for loans over \$1200:

$$\{t \mid t \in borrow \wedge t[amount] > 1200\}$$

This gives us all attributes, but suppose we only want the customer names. (We would use **project** in the algebra.)

We need to write an expression for a relation on scheme ($cname$).

$$\{t \mid \exists s \in borrow (t[cname] = s[cname] \wedge s[amount] > 1200)\}$$

In English, we may read this equation as "the set of all tuples t such that there exists a tuple s in the relation *borrow* for which the values of t and s for the *cname* attribute are equal, and the value of s for the *amount* attribute is greater than 1200."

The notation $\exists t \in r(Q(t))$ means "there exists a tuple t in relation r such that predicate $Q(t)$ is true".

How did we get the above expression? We needed tuples on scheme *cname* such that there were tuples in *borrow* pertaining to that customer name with *amount* attribute > 1200 .

The tuples t get the scheme *cname* implicitly as that is the only attribute t is mentioned with.

Let's look at a more complex example.

Find all customers having a loan from the SFU branch, and the cities in which they live:

$$\{t \mid \exists s \in borrow (t[cname] = s[cname] \wedge s[bname] = \text{"SFU"} \\ \wedge \exists u \in customer (u[cname] = s[cname] \wedge t[ccity] = u[ccity]))\}$$

In English, we might read this as "the set of all $(cname, ccity)$ tuples for which *cname* is a borrower at the SFU branch, and *ccity* is the city of *cname*".

Tuple variable s ensures that the customer is a borrower at the SFU branch.

Tuple variable u is restricted to pertain to the same customer as s , and also ensures that *ccity* is the city of the customer.

The logical connectives \wedge (AND) and \vee (OR) are allowed, as well as \neg (negation).

We also use the existential quantifier \exists and the universal quantifier \forall .

Some more examples:

1. Find all customers having a loan, an account, or both at the SFU branch:

$$\{t \mid \exists s \in borrow (t[cname] = s[cname] \wedge s[bname] = \text{"SFU"}) \\ \vee \exists u \in deposit (t[cname] = u[cname] \wedge u[bname] = \text{"SFU"})\}$$

Note the use of the \vee connective.

As usual, set operations remove all duplicates.

2. Find all customers who have **both** a loan and an account at the SFU branch.

Solution: simply change the \forall connective in 1 to a \wedge .

3. Find customers who have an account, but **not** a loan at the SFU branch.

$$\{t \mid \exists u \in deposit(t[cname] = u[cname] \wedge u[bname] = "SFU") \\ \wedge \neg \exists s \in borrow(t[cname] = s[cname] \wedge s[bname] = "SFU")\}$$

4. Find all customers who have an account at **all** branches located in Brooklyn.
(We used **division** in relational algebra.)

For this example we will use implication, denoted by a pointing finger in the text, but by \Rightarrow here.

The formula $P \Rightarrow Q$ means P implies Q , or, if P is true, then Q must be true.

$$\{t \mid \forall u \in branch (u[bcity] = "Brooklyn" \Rightarrow \\ \exists s \in deposit(t[cname] = s[cname] \wedge u[bname] = s[bname]))\}$$

In English: the set of all *cname* tuples *t* such that for all tuples *u* in the *branch* relation, if the value of *u* on attribute *bcity* is Brooklyn, then the customer has an account at the branch whose name appears in the *bname* attribute of *u*.

Division is difficult to understand. Think it through carefully.

Formal Definitions

1. A tuple relational calculus expression is of the form

$$\{t \mid P(t)\}$$

where P is a **formula**.

Several tuple variables may appear in a formula.

2. A tuple variable is said to be a **free variable** unless it is quantified by a \exists or a \forall . Then it is said to be a **bound variable**.
3. A formula is built of **atoms**. An atom is one of the following forms:
 - o $s \in r$, where *s* is a tuple variable, and *r* is a relation (\neq is not allowed).
 - o $s[x] \Theta u[y]$, where *s* and *u* are tuple variables, and *x* and *y* are attributes, and Θ is a comparison operator ($<, \leq, =, \neq, >, \geq$).
 - o $s[x] \Theta c$, where *c* is a constant in the domain of attribute *x*.
4. **Formulae** are built up from atoms using the following rules:
 - o An atom is a formula.

- o If P is a formula, then so are $\neg P$ and (P) .
- o If P_1 and P_2 are formulae, then so are $P_1 \vee P_2$, $P_1 \wedge P_2$ and $P_1 \Rightarrow P_2$.
- o If $P(s)$ is a formula containing a free tuple variable s , then

$$\exists s \in r(P(s)) \text{ and } \forall s \in r(P(s))$$

are formulae also.

5. Note some equivalences:
 - o $P_1 \wedge P_2 = \neg(\neg P_1 \vee \neg P_2)$
 - o $\forall t \in r(P(t)) = \neg \exists t \in r(\neg P(t))$
 - o $P_1 \Rightarrow P_2 = \neg P_1 \vee P_2$

Safety of Expressions

1. A tuple relational calculus expression may generate an infinite expression, e.g.

$$\{t \mid \neg(t \in borrow)\}$$

2. There are an infinite number of tuples that are not in *borrow*! Most of these tuples contain values that do not appear in the database.
3. **Safe Tuple Expressions**

We need to restrict the relational calculus a bit.

- o The domain of a formula P , denoted $\text{dom}(P)$, is the set of all values referenced in P .
- o These include values mentioned in P as well as values that appear in a tuple of a relation mentioned in P .
- o So, the domain of P is the set of all values explicitly appearing in P or that appear in relations mentioned in P .
- o $\text{dom}(t \in borrow \wedge t[amount] < 1200)$ is the set of all values appearing in *borrow*.
- o $\text{dom}(t \mid \neg(t \in borrow))$ is the set of all values appearing in *borrow*.

We may say an expression $\{t \mid P(t)\}$ is **safe** if all values that appear in the **result** are values from $\text{dom}(P)$.

4. A **safe** expression yields a finite number of tuples as its result. Otherwise, it is called **unsafe**.

Expressive Power of Languages

1. The tuple relational calculus restricted to safe expressions is equivalent in expressive power to the relational algebra

The Domain Relational Calculus

1. Domain variables take on values from an attribute's domain, rather than values for an entire tuple.

Formal Definitions

1. An expression is of the form

$$\{ \langle x_1, x_2, \dots, x_n \rangle \mid P(x_1, x_2, \dots, x_n) \}$$

where the $x_i, 1 \leq i \leq n$, represent domain variables, and P is a **formula**.

2. An atom in the domain relational calculus is of the following forms
 - o $\langle x_1, \dots, x_n \rangle \in r$ where r is a relation on n attributes, and $x_i, 1 \leq i \leq n$, are domain variables or constants.
 - o $x \Theta y$, where x and y are domain variables, and Θ is a comparison operator.
 - o $x \Theta c$, where c is a constant.
3. **Formulae** are built up from atoms using the following rules:
 - o An atom is a formula.
 - o If P is a formula, then so are $\neg P$ and (P) .
 - o If P_1 and P_2 are formulae, then so are $P_1 \vee P_2, P_1 \wedge P_2$ and $P_1 \Rightarrow P_2$.
 - o If $P(x)$ is a formula where x is a domain variable, then so are $\exists x(P(x))$ and $\forall x(P(x))$.

Example Queries

1. Find branch name, loan number, customer name and amount for loans of over \$1200.

$$\{ \langle b, l, c, a \rangle \mid \langle b, l, c, a \rangle \in \text{borrow} \wedge a > 1200 \}$$

2. Find all customers who have a loan for an amount > than \$1200.

$$\{ \langle c \rangle \mid \exists b, l, a (\langle b, l, c, a \rangle \in \text{borrow} \wedge a > 1200) \}$$

- Find all customers having a loan from the SFU branch, and the city in which they live.

$$\{ \langle c, x \rangle \mid \exists b, l, a (\langle b, l, c, a \rangle \in borrow \wedge b = \text{"SFU"} \wedge \exists y (\langle c, y, x \rangle \in customer)) \}$$

- Find all customers having a loan, an account or both at the SFU branch.

$$\{ \langle c \rangle \mid \exists b, l, a (\langle b, l, c, a \rangle \in borrow \wedge b = \text{"SFU"}) \vee \exists b, a, n (\langle b, a, c, n \rangle \in deposit \wedge b = \text{"SFU"}) \}$$

- Find all customers who have an account at **all** branches located in Brooklyn.

$$\{ \langle c \rangle \mid \forall x, y, z (\neg (\langle x, y, z \rangle \in branch) \vee z \neq \text{"Brooklyn"} \vee (\exists a, n (\langle x, a, c, n \rangle \in deposit))) \}$$

If you find this example difficult to understand, try rewriting this expression using implication, as in the tuple relational calculus example. Here's my attempt:

$$\{ \langle cn \rangle \mid \forall bn, as, bc ((\langle bn, as, bc \rangle \in branch \wedge bc = \text{"Brooklyn"}) \Rightarrow \exists an, ba (\langle cn, an, ba, bn \rangle \in deposit)) \}$$

I've used two letter variable names to get away from the problem of having to remember what x stands for.

Safety of Expressions

- As in the tuple relational calculus, it is possible to generate infinite expressions. The solution is similar for domain relational calculus-restrict the form to safe expressions involving values in the **domain** of the formula.

Read the text for a complete explanation.

Expressive Power of Languages

- All three of the following are equivalent:
 - The relational algebra.
 - The tuple relational calculus restricted to safe expressions.
 - The domain relational calculus restricted to safe expressions.

Modifying the Database

1. Up until now, we have looked at extracting information from the database. We also need to add, remove and change information. Modifications are expressed using the assignment operator.

Deletion

1. **Deletion** is expressed in much the same way as a query. Instead of displaying, the selected tuples are removed from the database. We can only delete whole tuples.

In relational algebra, a deletion is of the form

$$r \leftarrow r - E$$

where r is a relation and E is a relational algebra query.

Tuples in r for which E is true are deleted.

2. Some examples:

1. Delete all of Smith's account records.

$$deposit \leftarrow deposit - \sigma_{cname = "Smith"}(deposit)$$

2. Delete all loans with loan numbers between 1300 and 1500.

$$deposit \leftarrow deposit - \sigma_{loan\# \geq 1300 \wedge loan\# \leq 1500}(deposit)$$

3. Delete all accounts at Branches located in Needham.

$$r_1 \leftarrow \sigma_{bcity = "Needham"}(deposit \bowtie branch)$$

$$r_2 \leftarrow \Pi_{bname, account\#, cname, balance}(r_1)$$

$$deposit \leftarrow deposit - r_2$$

Insertions

1. To insert data into a relation, we either specify a tuple, or write a query whose result is the set of tuples to be inserted. Attribute values for inserted tuples must be members of the attribute's domain.
2. An insertion is expressed by

$$r \leftarrow r \cup E$$

where r is a relation and E is a relational algebra expression.

3. Some examples:

1. To insert a tuple for Smith who has \$1200 in account 9372 at the SFU branch.

$$deposit \leftarrow deposit \cup \{("SFU", 9372, "Smith", 1200)\}$$

2. To provide all loan customers in the SFU branch with a \$200 savings account.

$$r_1 \leftarrow (\sigma_{branch="SFU"}(borrow))$$

$$r_2 \leftarrow \Pi_{branch, loan\#, cname}(r_1)$$

$$deposit \leftarrow deposit \cup (r_2 \times \{(200)\})$$

Updating

1. Updating allows us to change some values in a tuple without necessarily changing all.

We use the update operator, δ , with the form

$$\delta_{A \leftarrow E}(r)$$

where r is a relation with attribute A , which is assigned the value of expression E .

The expression E is any arithmetic expression involving constants and attributes in relation r .

Some examples:

1. To increase all balances by 5 percent.

$$\delta_{balance \leftarrow balance * 1.05}(deposit)$$

This statement is applied to every tuple in *deposit*.

2. To make two different rates of interest payment, depending on balance amount:

$$\delta_{balance \leftarrow balance * 1.05}(\sigma_{balance > 10000}(deposit))$$

$$\delta_{balance \leftarrow balance * 1.06}(\sigma_{balance \leq 10000}(deposit))$$

Note: in this example the order of the two operations is important. (Why?)

Views

1. We have assumed up to now that the relations we are given are the actual relations stored in the database.
2. For security and convenience reasons, we may wish to create a personalized collection of relations for a user.
3. We use the term **view** to refer to any relation, not part of the conceptual model, that is made visible to the user as a "virtual relation".
4. As relations may be modified by deletions, insertions and updates, it is generally not possible to store views. (Why?) Views must then be recomputed for each query referring to them.

View Definition

1. A view is defined using the **create view** command:

create view v as <query expression>

where <query expression> is any legal query expression.

The view created is given the name *v*.

2. To create a view *all-customer* of all branches and their customers:

create view all-customer as
 $\Pi_{branch, names}(deposit) \cup \Pi_{branch, names}(borrow)$

3. Having defined a view, we can now use it to refer to the virtual relation it creates. View names can appear anywhere a relation name can.
4. We can now find all customers of the SFU branch by writing

$\Pi_{names}(\sigma_{branch = 'SFU'}(all - customer))$

Updates Through Views and Null Values

1. Updates, insertions and deletions using views can cause problems. The modifications on a view must be transformed to modifications of the actual relations in the conceptual model of the database.
2. An example will illustrate: consider a clerk who needs to see all information in the *borrow* relation except *amount*.

Let the view *loan-info* be given to the clerk:

```
create view loan-info as
   $\Pi_{branch, loan\#, names}(borrow)$ 
```

3. Since SQL allows a view name to appear anywhere a relation name may appear, the clerk can write:

```
loan-info  $\leftarrow$  loan-info  $\cup$  {( "SFU", 3, "Ruth" )}
```

This insertion is represented by an insertion into the actual relation *borrow*, from which the view is constructed.

However, we have no value for *amount*. A suitable response would be

- o Reject the insertion and inform the user.
- o Insert (`SFU`, 3, `Ruth`, null) into the relation.

The symbol **null** represents a null or place-holder value. It says the value is unknown or does not exist.

4. Another problem with modification through views: consider the view

```
create view branch-city as
   $\Pi_{branch, city}(borrow \bowtie customer)$ 
```

This view lists the cities in which the borrowers of each branch live.

Now consider the insertion

```
branch-city  $\leftarrow$  branch-city  $\cup$  {( "Brighton", "Woodside" )}
```

Using nulls is the only possible way to do this (see Figure 3.22 in the textbook).

If we do this insertion with nulls, now consider the expression the view actually corresponds to:

$$\Pi_{branch, city}(borrow \bowtie customer)$$

As comparisons involving nulls are always **false**, this query misses the inserted tuple.

To understand why, think about the tuples that got inserted into *borrow* and *customer*. Then think about how the view is recomputed for the above query