

# CS143 Notes: Relational Algebra

## Book Chapters

- (4th) Chapters 3.2
- (5th) Chapters 2.2-3
- (6th) Chapter 2.6

## Things to Learn

- Relational algebra
  - Select, Project, Join, ...

## Steps in Database Construction

1. Domain Analysis
2. Database design
3. Table creation: DDL
4. Load: bulk-load
5. Query and update: DML

## Database query language

### What is a query?

- Database jargon for question (complex word for simple concept)
- Questions to get answers from a database
  - Example: Get the students who are taking all CS classes but no Physics class
- Some queries are easy to pose, some are not
- Some queries are easy for DBMS to answer, some are not

## Relational query languages

- Formal: Relational algebra, relational calculus, datalog
- Practical: SQL ( $\leftarrow$  relational algebra), Quel ( $\leftarrow$  relational calculus), QBE ( $\leftarrow$  datalog)
- Relational Query:
  - Data sits in a disk
  - Submit a query
  - Get an answer

$$\text{Input relations} \longrightarrow \boxed{\text{query}} \longrightarrow \text{Output relation}$$

Executed against a set of relations and produces a relation

- \* Important to know
- \* *Very* useful: “Piping” is possible

## Relational Algebra

$$\text{Input relations (set)} \longrightarrow \boxed{\text{query}} \longrightarrow \text{Output relation (set)}$$

- Set semantics. no duplicate tuples. duplicates are eliminated
- In contrast, multiset semantics for SQL (performance reason)

## Examples to Use

- School information
  - Student(sid, name, addr, age, GPA)
- Class(dept, cnum, sec, unit, title, instructor)

sid	name	addr	age	GPA
301	John	183 Westwood	19	2.1
303	Elaine	301 Wilshire	17	3.9
401	James	183 Westwood	17	3.5
208	Esther	421 Wilshire	20	3.1

dept	cnum	sec	unit	title	instructor
CS	112	01	03	Modeling	Dick Muntz
CS	143	01	04	DB Systems	Carlo Zaniolo
EE	143	01	03	Signal	Dick Muntz
ME	183	02	05	Mechanics	Susan Tracey

- Enroll(sid, dept, cnum, sec)

sid	dept	cnum	sec
301	CS	112	01
301	CS	143	01
303	EE	143	01
303	CS	112	01
401	CS	112	01

### Simplest query: relation name

- **Query 1:** All students

### SELECT operator

Select all tuples satisfying a condition

- **Query 2:** Students with age < 18
- **Query 3:** Students with GPA > 3.7 and age < 18
- **Notation:**  $\sigma_C(R)$ 
  - Filters out rows in a relation
  - $C$ : A boolean expression with attribute names, constants, comparisons ( $>, \leq, \neq, \dots$ ) and connectives ( $\wedge, \vee, \neg$ )
  - $R$  can be either a relation or a result from another operator

### PROJECT operator

- **Query 4:** sid and GPA of all students
- **Query 5:** All departments offering classes
  - Relational algebra removes duplicates (set semantics)

- SQL does not (multiset or bag semantics)
- **Notation:**  $\pi_A(R)$ 
  - Filters out columns in a relation
  - A: a set of attributes to keep
- **Query 6:** sid and GPA of all students with age < 18
  - We can “compose” multiple operators
- **Q:** Is it ever useful to compose two projection operators next to each other?
- **Q:** Is it ever useful to compose two selection operators next to each other?

## CROSS PRODUCT (CARTESIAN PRODUCT) operator

- Example:  $R \times S$

A	B
a <sub>1</sub>	b <sub>1</sub>
a <sub>2</sub>	b <sub>2</sub>
a <sub>2</sub>	b <sub>3</sub>

 $\times$ 

A	B
a <sub>1</sub>	b <sub>1</sub>
a <sub>1</sub>	b <sub>2</sub>
a <sub>1</sub>	b <sub>3</sub>
a <sub>2</sub>	b <sub>1</sub>
a <sub>2</sub>	b <sub>2</sub>
a <sub>2</sub>	b <sub>3</sub>

 $=$ 

A	B
a <sub>1</sub>	b <sub>1</sub>
a <sub>1</sub>	b <sub>2</sub>
a <sub>1</sub>	b <sub>3</sub>
a <sub>2</sub>	b <sub>1</sub>
a <sub>2</sub>	b <sub>2</sub>
a <sub>2</sub>	b <sub>3</sub>

- Concatenation of tuples from both relations
- One result tuple for each pair of tuples in  $R$  and  $S$
- If column names conflict, prefix with the table name
- **Notation:**  $R_1 \times R_2$ 
  - $R_1 \times R_2 = \{t \mid t = \langle t_1, t_2 \rangle \text{ for } t_1 \in R_1 \text{ and } t_2 \in R_2\}$
- **Q:** Looks odd to concatenate unrelated tuples. Why use  $\times$ ?
- **Query 7:** Names and addresses of students who take CS courses with GPA < 3

- **Q:** Can we write it differently?
  - Benefit of RDBMS. It figures out the best way to compute.
- **Q:** If  $|R| = r$  and  $|S| = s$ , what is  $|R \times S|$ ?

## NATURAL JOIN operator

- Example: Student  $\bowtie$  Enroll
  - Shorthand for  $\sigma_{Student.sid=Enroll.sid} (Student \times Enroll)$
- **Notation:**  $R_1 \bowtie R_2$ 
  - Concatenate tuples horizontally
  - Enforce equality on common attributes
  - We may assume only one copy of the common attributes are kept
- **Query 8:** Names of students taking CS classes
  - Explanation: start with the query requiring sid, not name
- **Query 9:** Names of students taking classes offered by “Carlo Zaniolo”

- Natural join: The most natural way to join two tables

## THETA JOIN operator

- **Example:** Student  $\bowtie_{Student.sid=Enroll.sid \wedge GPA > 3.7}$  Enroll
- **Notation:**  $R_1 \bowtie_C R_2 = \sigma_C(R_1 \times R_2)$
- Generalization of natural join

- Often implemented as the basic operation in DBMS

## RENAME operator

- **Query 10:** Find the pairs of student names who live in the same address.
- What about  $\pi_{name,name}(\sigma_{addr=addr}(\text{Student} \times \text{Student}))$ ?
- **Notation:**  $\rho_S(R)$  – rename  $R$  to  $S$
- **Notation:**  $\rho_{S(A1',A2')}(R)$  for  $R(A1, A2)$  – rename  $R(A1, A2)$  to  $S(A1', A2')$
- **Q:** Is  $\pi_{Student.name,S.name}(\sigma_{Student.addr=S.addr}(\text{Student} \times \rho_S(\text{Student})))$  really correct?
  - How many times (John, James) returned?

## UNION operator

- **Query 11:** Find all student and instructor names.
  - **Q:** Can we do it with cross product or join?
- **Notation:**  $R \cup S$ 
  - Union of tuples from  $R$  and  $S$
  - The schemas of  $R$  and  $S$  should be the same
  - No duplicate tuples in the result

## DIFFERENCE operator

- **Query 12:** Find the courses (dept, cnum, sec) that no student is taking
  - How can we find the courses that at least one student is taking?
- **Notation:**  $R - S$ 
  - Schemas of  $R$  and  $S$  must match exactly

- **Query 13:** What if we want to get the titles of the courses?

– Very common. To match schemas, we lose information. We have to join back.

## INTERSECT operator

- **Query 14:** Find the instructors who teach both CS and EE courses
  - **Q:** Can we answer this using only selection and projection?

- **Notation:**  $R \cap S = R - (R - S)$ 
  - Draw Venn Diagram to verify

## DIVISION operator

Use the boards very carefully keeping all examples.

- Division operator is not used directly by any one
- But how we compute the answer for division is very important to learn
- Learn how we computed the answer, not the operator
- **Query 15:** Find the student sids who take every CS class available
  - **Q:** What will be the answer?

– **Q:** How can we compute it?  
*Give time to think about*

- \* **Step 1:** We need to know which student is taking which class. Where do we get the student sid and the courses they take(sid, dept, cnum, sec)?

\* **Step 2:** We also need to know all CS classes. How do we get the set of all CS courses (dept, cnum, sec)?

\* **Step 3:** What does the relation from Step 1 look like if all students take all CS courses?

How can we compute this? Let us call this relation  $R$

\* **Step 4:** What does  $(R - \text{Enroll})$  look like? What's its meaning?

\* **Step 5:** What is the meaning of the projection of Step 4?

\* **Step 6:** How can we get the student sids who take all CS courses?

• **Notation:**  $R/S$ .

– **Query 16:** Find All  $A$  values in  $R$  such that the values appear with *all*  $B$  values in  $S$ .

$R$		$S$
$A$	$B$	$B$
$a_1$	$b_1$	$b_1$
$a_1$	$b_2$	$b_2$
$a_2$	$b_1$	
$a_3$	$b_2$	

$R \approx$  *students and classes they take*

$S \approx$  *all CS classes*  $R/S \approx$  *All students who take all CS classes*

– Formal definition:

\* We assume  $R(A, B)$  and  $S(B)$

\* The set of all  $a \in R.A$  such that  $\langle a, b \rangle \in R$  for *every*  $b \in S$

\*  $R/S = \{a \mid a \in R.A \text{ and } \langle a, b \rangle \in R \text{ for all } b \in S\}$

– Result for the example:

$A$
$a_1$

– **Q:** How to compute it?



*Ans:*  $R/S = \pi_A(R) - \pi_A((\pi_A(R) \times S) - R)$

All  $(R.A, S.B)$  pairs

- \*  $\pi_A(R)$ : All  $A$ 's

- \*  $S$ : All  $B$ 's

- \*  $\pi_A(R) \times S$ : all  $R.A$  and  $S.B$  pairs

$\pi_A(R) \times S - R$ :  $\langle A, B \rangle$  pairs that are missing in  $R$

$\pi_A(\pi_A(R) \times S - R)$ : All  $R.A$ 's that do not have some  $S.B$

- Analogy with interger division

- \* In integer:  $R/S$  is the largest integer  $T$  such that  $T \times S \leq R$

- \* In relational algebra:  $R/S$  is the largest relation  $T$  such that  $T \times S \subseteq R$

- The division operator is not used often, but how to compute it is important

## More questions

- **Q:** sids of students who did not take any CS courses?

- **Q:** Is  $\pi_{sid}(\sigma_{title \neq 'CS'}(Enroll))$  correct?

- **Q:** What is its complement?

- **General advice:** When a query is difficult to write, think in terms of its complement.

## Relational algebra: things to remember

- Data manipulation language (query language)

- Relation  $\rightarrow$  algebra  $\rightarrow$  relation

- Relational algebra: set semantics, SQL: bag semantics

- Operators:  $\sigma, \times, \bowtie, \rho, \cup, -, \cap, /$

- General suggestion: If difficult to write, consider its complement