

Unit-4 (Contd..)

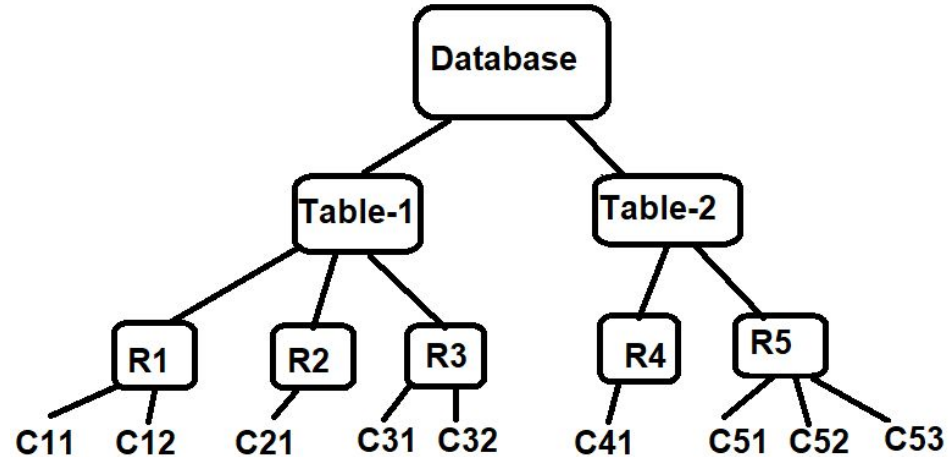
- Multiple Granularity of locks
 - Granularity hierarchy
 - Intention Lock modes
 - Compatibility matrix for lock modes
- Handling Deadlocks
 - Deadlock Prevention
 - Deadlock Detection
 - Deadlock Recovery
- Multiversion schemes
 - Multiversion using Timestamp ordering
 - Multiversion using 2-phase locking

Situation when transaction holds multiple data item

- A Transaction T_i access more than one data item from database.
- Data item may be a column's value, row, table or entire database
- Database \rightarrow Tables \rightarrow Rows \rightarrow Columns
- If a Transaction needs to access a database:
 - Should it needs to lock entire database? If so, it needs to apply multiple locks on its tables, rows and columns
- If a Transaction needs to access a specific data item (Q):
 - Should it needs to lock only that data item (Q) or entire database?
- Both in the above cases, the Intensity of locks differ

Granularity hierarchy

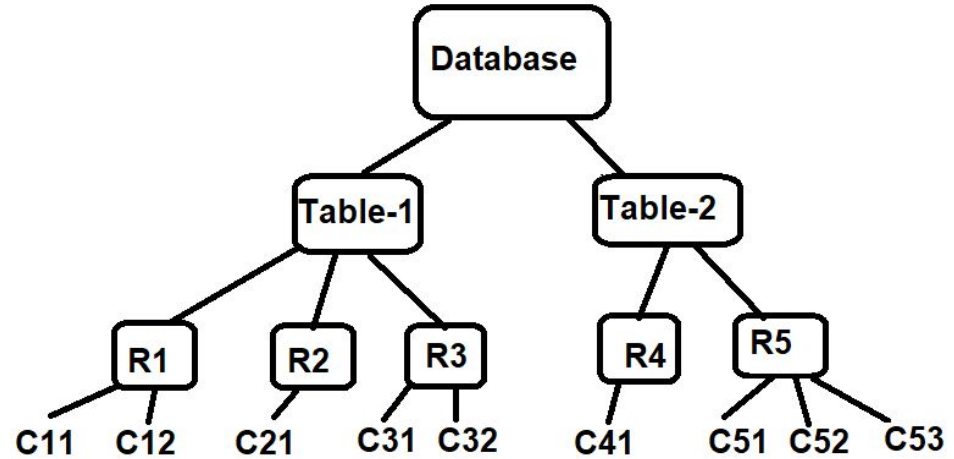
- Represents the level/effect of locks on data items requested by transaction.
- A database is a collection of tables, rows and columns grouped in a hierarchical manner.
- All Data items are grouped under sub-tree nodes and sub-tree nodes are grouped under root node.
- Two locks can be applied by default:
 - Shared lock – For read operations
 - Exclusive lock – for both read and write operations



Note: R means Row , C means Column

Granularity hierarchy

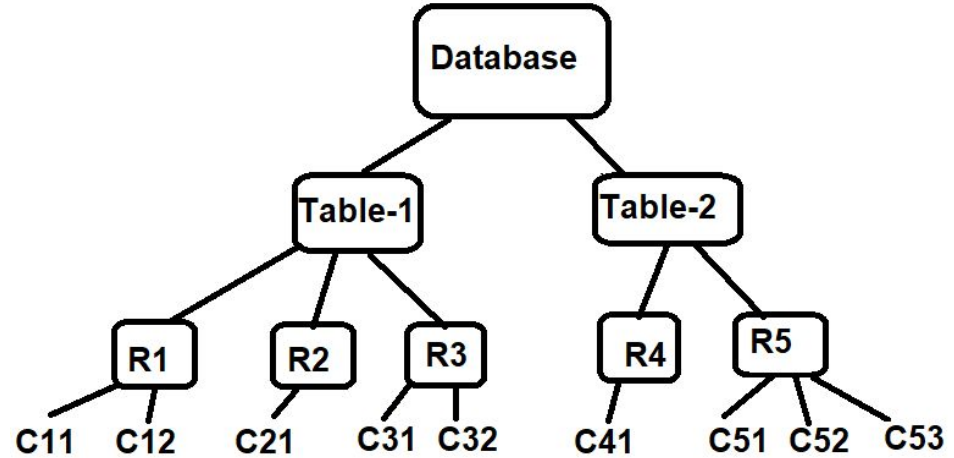
- A Lock can be explicit lock or implicit lock.
- Example: If Ti explicitly call shared lock on R5 then its descendants automatically will be applied with implicit locks.
- If Tj request X-lock on C51 explicitly then it will be rejected due to an implicit already applied before by Ti.
- If Tk request a lock on Database, it will be rejected as its child node R5 is locked by Ti.



Note: R means Row , C means Column

Granularity hierarchy

- How Transactions T_i , T_j or T_k determines that other transactions are holding locks on data items?
 - They should traverse from root node to current node and check the lock is compatible or not with other transactions.
 - Obviously it's a time consuming process
- To overcome this, we use three more locks called '**Intention lock modes**'



Note: R means Row , C means Column

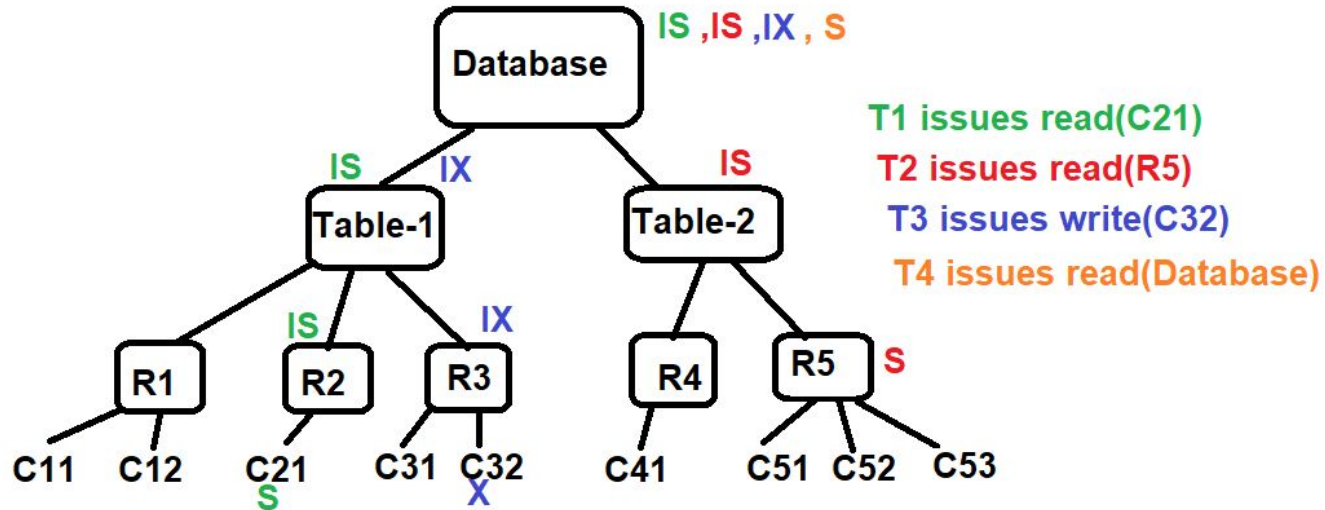
Granularity Locks

- **Shared lock (S)** – to read data item
- **Exclusive lock (X)** – to read/write data item
- **Intention Shared lock (IS)** – If Node N is Shared locked, then all its ancestors are locked in IS lock
- **Intention Exclusive lock (IX)** – If Node N is Exclusive locked, then all its ancestors are locked in IX lock
- **Shared Intention Exclusive lock (SIX)** – If Node N is Shared locked, its sub-Node is Exclusive locked then ancestors of N are locked in SIX mode lock

Lock Compatibility matrix for Multiple Granularity of locks

	IS	IX	S	SIX	X
IS	true	true	true	true	false
IX	true	true	false	false	false
S	true	false	true	false	false
SIX	true	false	false	false	false
X	false	false	false	false	false

Example



Note: R means Row , C means Column

which of them run concurrent?

T1, T2

T1, T3

T3, T4

T1, T2, T4

Multiple Granularity Locking Scheme

- Transaction T_i can lock a node Q , using the following rules:
 1. The lock compatibility matrix must be observed.
 2. The root of the tree must be locked first, and may be locked in any mode.
 3. A node Q can be locked by T_i in S or IS mode only if the parent of Q is currently locked by T_i in either IX or IS mode.
 4. A node Q can be locked by T_i in X, SIX, or IX mode only if the parent of Q is currently locked by T_i in either IX or SIX mode.
 5. T_i can lock a node only if it has not previously unlocked any node (that is, T_i is two-phase).
 6. T_i can unlock a node Q only if none of the children of Q are currently locked by T_i .
- Observe that locks are acquired in root-to-leaf order, whereas they are released in leaf-to-root order.

Handling Deadlocks

- What is Deadlock – A situation in which two transactions T_i , T_j mutually depends on each other to access data item which will be a never ending transaction.
- Example: $T_1 \rightarrow T_2 \rightarrow T_3 \rightarrow T_4 \rightarrow T_1$
- Some transactions facing Deadlocks can be recoverable but some cannot.
- Deadlock Prevention
- Deadlock Detection
- Deadlock Recovery

T_3	T_4
lock-X(B) read(B) $B := B - 50$ write(B)	lock-S(A) read(A) lock-S(B)
lock-X(A)	

Deadlock prevention

Deadlock prevention protocols ensure that the system will never enter into a deadlock state.

Using Timestamp ordering mechanism of transactions in order to predetermine a deadlock situation

a) Wait-Die Scheme -- allows older transaction to wait and younger one to die.

If a transaction T_i requests to lock a data item, which is already held with a conflicting lock by another transaction T_j , then one of the two possibilities may occur:

- If $TS(T_i) < TS(T_j)$ then T_i is allowed to wait till data item is available.
- If $TS(T_i) > TS(T_j)$ then T_i dies and would be restarted later

b) Wound-Wait Scheme -- allows the younger transaction to wait; but when an older transaction requests an item held by a younger one, the older transaction forces the younger one to abort and release the item.

- If $TS(T_i) < TS(T_j)$ then T_i forces T_j to rollback and T_i waits till data item available. i.e T_i wounds T_j .
- If $TS(T_i) > TS(T_j)$ then T_i is forced to wait till data item is available.

Deadlock prevention


- **Timeout-Based Schemes:**

- A transaction waits for a lock only for a specified amount of time. After that, the wait times out and the transaction is rolled back.
- Ensures that deadlocks get resolved by timeout if they occur
- Simple to implement
- But may roll back transaction unnecessarily in absence of deadlock
- Difficult to determine good value of the timeout interval.
- Starvation is also possible


Deadlock Detection

- **Wait-for graph**

- *Vertices:* transactions
- *Edge from $T_i \rightarrow T_j$:* if T_i is waiting for a lock held in conflicting mode by T_j
- The system is in a deadlock state if and only if the wait-for graph has a cycle.
- Invoke a deadlock-detection algorithm periodically to look for cycles.



Wait-for graph without a cycle



Wait-for graph with a cycle

Deadlock Recovery

- When deadlock is detected :
 - Some transaction will have to be rolled back (made a **victim**) to break deadlock cycle.
 - Using **Victim Selection**, Select a transaction as victim that will have minimum risk
 - Rollback -- determine how far to roll back transaction
 - **Total rollback**: Abort the transaction and then restart it.
 - **Partial rollback**: Roll back victim transaction only as far as necessary to release locks that another transaction in cycle is waiting for

Multi version Schemes

- It is a mechanism of creating a new copy of data item when a write operation is performed by a transaction on Data Item-Q.
- The new copy is called as Version and each version is identified as $\{Q_1, Q_2, Q_3 \dots Q_n\}$ where Q_1, Q_n are the older, newer versions of data item Q respectively.
- Idea behind this scheme:
 - To allow a transaction T_i to perform read operation on Q while other Transaction T_j performs conflicting write operation on same data item Q concurrently.

Multi version Schemes

- Multiversion schemes keep old versions of data item to increase concurrency. Several variants:
 - **Multiversion Timestamp Ordering**
 - **Multiversion Two-Phase Locking**
- Key ideas:
 - Each successful **write** results in the creation of a new version of the data item written.
 - Use timestamps to label versions.
 - When a **read**(Q) operation is issued, select an appropriate version of Q based on the timestamp of the transaction issuing the read request, and return the value of the selected version.
 - **reads** never have to wait as an appropriate version is returned immediately.

Multi version schemes

Each data item Q has a sequence of versions $\langle Q_1, Q_2, \dots, Q_m \rangle$.

Each version Q_k contains three data fields:

Content -- the value of version Q_k

W-timestamp(Q_k) -- timestamp of the transaction that created
(wrote) version Q_k

R-timestamp(Q_k) -- largest timestamp of a transaction that
successfully read version Q_k

Data Item - Q		
Content	W-Timestamp	R-Timestamp

Multi version using Timestamp Ordering

- Suppose that transaction T_i issues a **read**(Q) or **write**(Q) operation. Let Q_k denote the version of Q whose write timestamp is the largest write timestamp less than or equal to $TS(T_i)$.
 1. If transaction T_i issues a **read**(Q), then
 - the value returned is the content of version Q_k
 - If $R\text{-timestamp}(Q_k) < TS(T_i)$, set $R\text{-timestamp}(Q_k) = TS(T_i)$,
 2. If transaction T_i issues a **write**(Q)
 1. if $TS(T_i) < R\text{-timestamp}(Q_k)$, then transaction T_i is rolled back.
 2. if $TS(T_i) = W\text{-timestamp}(Q_k)$, the contents of Q_k are overwritten
 3. Otherwise, a new version Q_i of Q is created
 - $W\text{-timestamp}(Q_i)$ and $R\text{-timestamp}(Q_i)$ are initialized to $TS(T_i)$.

Multiversion Two-phase locking

- Differentiates between read-only transactions and update transactions
- **Update transactions** acquire read and write locks, and hold all locks up to the end of the transaction. That is, update transactions follow rigorous two-phase locking.
 - Read of a data item returns the latest version of the item
 - The first **write** of Q by T_i results in the creation of a new version Q_i of the data item Q written
 - $W\text{-timestamp}(Q_i)$ set to ∞ initially
 - When update transaction T_i completes, commit processing occurs:
 - Value **ts-counter** stored in the database is used to assign timestamps
 - **ts-counter** is locked in two-phase manner
 - Set $TS(T_i) = \text{ts-counter} + 1$
 - Set $W\text{-timestamp}(Q_i) = TS(T_i)$ for all versions Q_i that it creates
 - **ts-counter** = **ts-counter** + 1

Multiversion Two-phase locking

- **Read-only transactions**

- are assigned a timestamp = **ts-counter** when they start execution
- follow the multiversion timestamp-ordering protocol for performing reads
 - Do not obtain any locks

- Read-only transactions that start after T_i increments **ts-counter** will see the values updated by T_i .
- Read-only transactions that start before T_i increments the **ts-counter** will see the value before the updates by T_i .
- Only serializable schedules are produced.

Multiversion Timestamp ordering - Example

Multiversion Timestamp ordering			step 0			Step 3 Read by T2		
	TST	RTS	Content	WTS	RTS	If $RTS < TST$ Update $RTS = TST$		
	5	10	X0	0	0	Content	WTS	RTS
Step	T1	T2				X1	5	10
1	Read X		Step 1 Read by T1					
2	Write X		If $RTS < TST$ Update $RTS = TST$			Step 4 Write by T2		
3		Read X	If $RTS > TST$ read older version			TST > WTS Create new version		
4		Write X	Content	WTS	RTS	Change both WTS and $RTS = TST$		
5	Read X		X0	0	5	Content	WTS	RTS
6	Write X					X2	10	10
7		Write X	Step 2 Write by T1					
			If $TST < RTS$ then Roll back			Step 5 Read by T1		
Step 6	Write by T1		If $TST = WTS$ Overwrite content			If $RTS > TST$ read older version		
Check latest version of X			If $TST > WTS$ Create new version			where $TST = RTS$		
Content	WTS	RTS	Change both WTS and $RTS = TST$			Do not update anything		
X2	10	10	Content	WTS	RTS	Content	WTS	RTS
TST < RTS so Roll Back			X1	5	5	X1	5	5
Since T1 had aborted T2 has to abort as it has read contents written by T1								
Remove version X0 and X1 as their WTS < least TST								
The multiversion timestamp-ordering scheme has the desirable property that a read request never fails and is never made to wait								
Undesirable properties								
Reading of a data item also requires the updating of the R-timestamp field, resulting in two potential disk accesses, rather than one.								
Second, the conflicts between transactions are resolved through rollbacks, rather than through waits.								